

Part 2: Clustering (45 points)

Unsupervised data mining methods are fairly diverse. At one extreme, there are the pure statistical summarization techniques that are the mainstay of exploratory data analysis, somewhere in between are data transformations like principal components analysis, which have a deeper mathematical underpinning but a fairly mechanical formulation, and at the other extreme you have unsupervised models that make strong assumptions about the data distribution. Clustering methods fall into that latter category: they make an assumption about instance similarity given attributes and use this modeling assumption to provide an interpretable division of a dataset.

The strong modeling assumptions made by clustering algorithms have varying benefits and drawbacks. In some cases, the modeling assumption of a clustering algorithm adheres to our intuitions about the data. For example, we may use a Gaussian mixture model for normally distributed data, and get extremely useful results. In other cases, the modeling assumptions are clearly incorrect, but these assumptions may provide a simpler formulation that pragmatically works quite well for a dataset. As an example, in mixture modeling for text corpora (like the news example in class), we make the assumption that documents are built by randomly drawing words independently of each other (but dependent on the cluster) -- this is not how journalists write news articles, but it turns out this simplifying assumption allows fairly powerful models. In the worst case, the modeling assumptions conflict with the underlying data generating process and impede our ability to understand the data. For example, in class we discussed how distance functions do not always work as we expect, and two people who both like food and music may have the same distance as two people who both like food, music, hiking, biking, camping, and rock climbing, even though we expect that the first pair will have far less to talk about than the second pair.

In this part of the assignment, we'll use the scikit-learn implementations of clustering (<http://scikit-learn.org/stable/modules/clustering.html>) in the clustering and mixture modules. We'll explore four different clustering methods (k-Means (<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>), Gaussian mixture models (<http://scikit-learn.org/stable/modules/mixture.html#gmm>), and Agglomerative clustering (<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.AgglomerativeClustering.html>)) and try to better understand the modeling assumptions each makes.

```

In [6]: ## Preliminaries

#Show plots in the notebook
%matplotlib inline

from sklearn import datasets, preprocessing, cross_validation, feature_
from sklearn import linear_model, svm, metrics, ensemble, neighbors, mi
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import urllib2
import random
import math

# Helper functions
def folds_to_split(data, targets, train, test):
    data_tr = pd.DataFrame(data).iloc[train]
    data_te = pd.DataFrame(data).iloc[test]
    labels_tr = pd.DataFrame(targets).iloc[train]
    labels_te = pd.DataFrame(targets).iloc[test]
    return [data_tr, data_te, labels_tr, labels_te]

def computeLabelPercentage(data, axis=0):
    return data.groupby(by=axis).size().apply(lambda x: x/data.count())

```

Question 1: Implementing the basic k-Means algorithm (15 points)

Let's work through the writing the code for a basic k-means implementation. Assume you have data this is either a numpy ndarray or pandas DataFrame.

1. Write a function `sample_centroids` that takes as parameters `k` and your data, then picks `k` (distinct) random data points as centroids. (The function `random.sample` may be helpful)
2. Write a function `find_closest_centroid` that takes as parameters a list of centroids and your data, then computes the distance of each instance to each of `k` centroids, and returns the minimum centroid for each instance. Use euclidean distance in your implementation.
3. Write a function `update_centroid` that takes as a parameter a dataset containing the instances mapped to a particular centroid and returns the new centroid
4. Put these together in a function `kmeans` that samples centroids and loops over the steps of finding the closest centroid for each instance and updating the centroids until no centroid assignments change.
5. Put it all together: generate some sample data using `make_blobs` and confirm that your k-Means implementation works. At each iteration, plot the data points and cluster assignments. (See the [scikit-learn clustering comparison \(http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html\)](http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html) example if you need some help with figuring out the plotting)

In [7]: *# 1. sample_centroids*

```
def sample_centroids(k, data):  
    data_clean = data  
    indexes = random.sample(data.index, k)  
    centroids = data.iloc[indexes]  
    return [centroids, indexes]
```

In [8]: *# 2. find_closest_centroid*

```
def find_closest_centroid(centroids, data):  
    data_values = data.values  
    inst_cent_map = []  
  
    for row in range(0, len(data_values)):  
        centroid_distances = []  
  
        for centroid in range(0, len(centroids)):  
            distance = 0  
            cent = centroids[centroid]  
  
            for col in range(0, len(data_values[0])):  
                x = cent[col]  
                y = data_values[row][col]  
                distance += (x - y)**2  
  
            distance = math.sqrt(distance)  
            centroid_distances.append(distance)  
            closest_centroid = centroid_distances.index(min(centroid_distances))  
            inst_cent_map.append([data_values[row], centroids[closest_centroid]])  
    return inst_cent_map
```

In [9]: # 3. update_centroid

```
def update_centroid(instances_centroid_map, k_means):
    new_centroids = []
    for i in range(0, k_means):
        same_centroid_instances = []

        added_initial_centroid = 0
        for instance in range(0, len(instances_centroid_map)):
            instance_features = instances_centroid_map[instance][0]
            centroid_points = instances_centroid_map[instance][1]
            centroid_group = instances_centroid_map[instance][2]

            if (centroid_group == i):
                if (added_initial_centroid != 1):
                    same_centroid_instances.append(centroid_points)
                    added_initial_centroid = 1
                    same_centroid_instances.append(instance_features)

        new_centroid = []
        if (len(same_centroid_instances) > 0):
            for feature in range(0, len(same_centroid_instances[0])):
                feature_total = 0
                total_instances = len(same_centroid_instances)

                for inst in range(0, total_instances):
                    instance = same_centroid_instances[inst]
                    feature_total += instance[feature]

                new_centroid.append(round(feature_total / total_instances))

            new_centroids.append([new_centroid, total_instances])
    return new_centroids
```

In [10]:

```

# 4. Generating data and putting it all together
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches

[data, targets] = datasets.make_blobs(n_samples=10000, n_features=2,
                                     random_state=20160217)

data = pd.DataFrame(data)
targets = pd.DataFrame(targets)

def k_means():
    # Sample centroids
    data_clean = data
    k_means = 2
    [centroids, indexes] = sample_centroids(k_means, data)

    # init total_instances
    total_instances = []
    for i in range(0, k_means):
        total_instances.append(1)

    data_clean.drop(data_clean.index[indexes], inplace=True)
    centroids = centroids.values

    centroid_updated = 1;
    while (centroid_updated == 1):

        centroid_updated = 0

        instances_centroids_map = find_closest_centroid(centroids, data

        old_centroids = centroids
        centroids_info = update_centroid(instances_centroids_map, k_me

        centroids = []

        new_total_instances = []
        for i in range(0, len(centroids_info)):
            centroids.append(centroids_info[i][0])
            new_total_instances.append(centroids_info[i][1])

        if total_instances == new_total_instances:
            centroid_updated = 0
            centroids = old_centroids
        else:
            centroid_updated = 1
            total_instances = new_total_instances

        x_points_red = []
        y_points_red = []

        x_points_blue = []
        y_points_blue = []

```

```

x_centroids = []
y_centroids = []

instances_centroids_map = find_closest_centroid(centroids, data)
for inst in range(0, len(instances_centroids_map)):
    instance = instances_centroids_map[inst][0]

    if (centroids[0] == instances_centroids_map[inst][1]):
        x_points_red.append(instance[0])
        y_points_red.append(instance[1])
    else:
        x_points_blue.append(instance[0])
        y_points_blue.append(instance[1])

for cent in range(0, len(centroids)):
    x_centroids.append(centroids[cent][0])
    y_centroids.append(centroids[cent][1])


plt.scatter(x_points_red, y_points_red, c='r', alpha=0.5)
plt.scatter(x_points_blue, y_points_blue, c='b', alpha=0.5)




plt.scatter(x_centroids, y_centroids, c='g', s=100)
plt.show()

return [centroids, data]

```

```
[centroids, data] = k_means()
```





Flower Arrangements with Python

In this assignment, we'll look at three different clustering algorithms. The first is k-Means -- you're pretty familiar with that algorithm since you just implemented it. Not that your code isn't beautiful, but we'll be using the [scikit-learn](http://scikit-learn.org/) K-Means ([http://scikit-](http://scikit-learn.org/)

learn.org/stable/modules/generated/sklearn.cluster.KMeans.html) to make sure we have consistent results

Next, is hierarchical clustering, also known as agglomerative clustering. Hierarchical clustering starts out with each instance as its own cluster, then merges the two closest clusters, thus reducing the number of clusters by 1. This merging continues until only the desired number of clusters remain. The main implementation question in hierarchical clustering is how to compute the distance between two clusters of points. The [sklearn implementation \(http://scikit-learn.org/stable/modules/generated/sklearn.cluster.AgglomerativeClustering.html\)](http://learn.org/stable/modules/generated/sklearn.cluster.AgglomerativeClustering.html) offers a few options.

- average: compute all pairwise distances between the two clusters, and then take the average of these distances
- complete (also known as maximum) linkage: compute all pairwise distances and use the largest distance
- Ward: compute the change in cluster variance after all possible cluster merges and choose the merge that results in the smallest variance increase

Finally, there are [Gaussian Mixture Models \(http://scikit-learn.org/stable/modules/mixture.html#gmm\)](http://scikit-learn.org/stable/modules/mixture.html#gmm). These models assume that the attribute values in each dimension follow a Gaussian (or normal) distribution. One nice feature of this assumption is that the Gaussian models can account for differing attribute variances. For example, if an "income" attribute has a very large range, and a "years of education" has a small range, the fit variance for the former attribute will be large, and for the latter attribute will be small. Of course, you already know how to achieve a somewhat similar result in other algorithms, right?

Let's look at three different clustering algorithms on the Iris dataset. In the plot below, the columns correspond to the three clustering algorithms and ground truth - the flower species. The rows correspond to a combination of two of the four attributes (petal length, petal width, sepal length, sepal width), of which there are six (4 choose 2). Given the four attributes, you'll have to think high-dimensionally to really get a sense of the clusterings. You can see that all of the clustering algorithms roughly correspond to the true "clusters" that are based on the the flower species. However, each behaves slightly differently. Hierarchical clustering with complete linkage will prefer small, compact clusters. You can see in the last row of the figure how the red cluster for hierarchical clustering remains small, while the other algorithms have more mixing. In that last row, you can also notice how the Gaussian mixture model estimates a higher variance for the attributes in the blue cluster (which have outliers in the top right), resulting in more cluster mixing. By comparison, the same cluster in k-Means has its center pulled towards the extreme which results in a smaller cluster focused on the outliers.

```

In [11]: iris = datasets.load_iris()
print pd.DataFrame(iris.data).describe()
clusterers = {
    'KMeans':cluster.KMeans(n_clusters=3, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=3, linkag
    'GMM':mixture.GMM(n_components=3, random_state=20160217),
    }

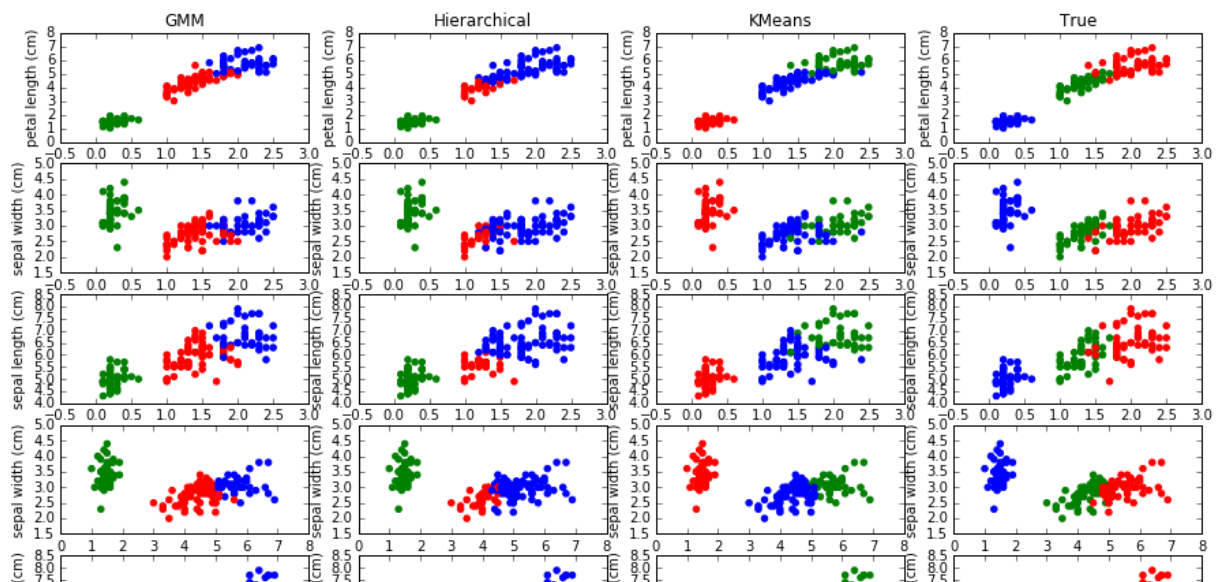
clusterings = {'True':iris.target}

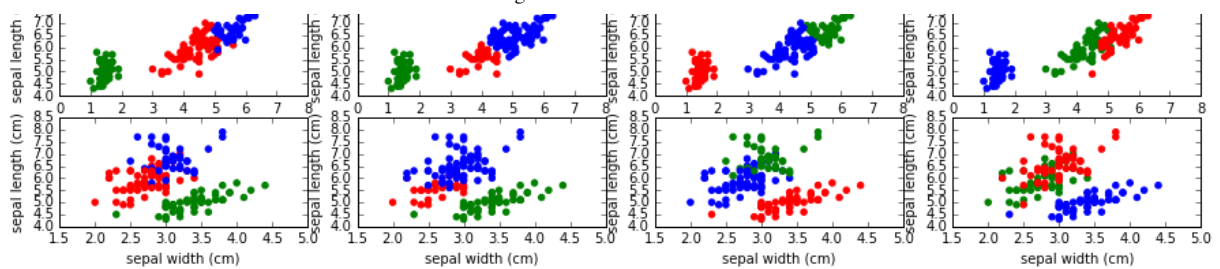
for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(iris.data)

colors = np.array([x for x in 'bgcrmykbgrcmykbgrcmykbgrcmyk'])
colors = np.hstack([colors] * 20)
plt.figure(figsize=(15,10))
pnum=1
for x in range(3,0,-1):
    for y in range(x-1, -1, -1):
        for clusterer in sorted(clusterings.keys()):
            plt.subplot(6,4,pnum)
            if pnum < 5: plt.title(clusterer)
            plt.xlabel(iris.feature_names[x])
            plt.ylabel(iris.feature_names[y])
            plt.scatter(iris.data[:,x], iris.data[:,y], color=colors[cl
            pnum+=1

```

	0	1	2	3
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000





```
In [12]: # Load Heart Data
heart_data = urllib2.urlopen("http://archive.ics.uci.edu/ml/machine-lea
heart = pd.read_csv(heart_data, quotechar='"', skipinitialspace=True, n
heart=heart.dropna()
heart_attrs = heart.ix[:, :-1]
heart_labels = heart.ix[:, -1]
heart_processed = pd.DataFrame(preprocessing.StandardScaler().fit_trans
heart_labels_values = heart_labels.values
```

Question 2: Clustering Heart Disease (15 points)

In this exercise, you'll be clustering the heart disease data we've seen in a couple of assignments already. I've loaded and processed the data for you. You'll be using the three clusterers from the example above (KMeans, Hierarchical, GMM), using the same `random_state` settings and initialization as in the example. Cluster the processed heart data using five clusters and the three clustering algorithms.

- Compute the cluster sizes. What do you notice about the cluster sizes for each algorithm?
- Compute the mean value of the HeartDisease column (`heart_labels`) for each cluster. How well do the clusters correspond to differing heart disease conditions?
- Write a function that `wss` computes the within-cluster sum of square errors (MSE of a cluster) and a function `bss` to compute the between-cluster sum of square errors (see the lecture slides!), and report both metrics for each of the clustering methods.
- For each of the clusterings (15 total), choose one attribute that has a different distribution than the overall data, and discuss how the distribution differs for that attribute.
- Make a figure showing boxplots of each attribute in each clustering, as well as a figure with a boxplot of all the data. (The figure should have three rows [one for each clustering algorithm] and six columns [one for all the data, five for each of the clusterings], and each boxplot will have twelve attributes.

Answers 2.

1. Cluster sizes

Clustering Method: Hierarchical

Cluster: 0 size: 131

Cluster: 1 size: 129

Cluster: 2 size: 13

Cluster: 3 size: 1

Cluster: 4 size: 23

Clustering Method: KMeans

Cluster: 0 size: 39

Cluster: 1 size: 65

Cluster: 2 size: 62

Cluster: 3 size: 46

Cluster: 4 size: 85

Clustering Method: GMM

Cluster: 0 size: 60

Cluster: 1 size: 89

Cluster: 2 size: 78

Cluster: 3 size: 43

Cluster: 4 size: 27

2. Cluster means

Clustering Method: Hierarchical

Cluster: 0 mean: 0.4275

Cluster: 1 mean: 1.1705

Cluster: 2 mean: 2.5385

Cluster: 3 mean: 0.0

Cluster: 4 mean: 1.7826

Clustering Method: KMeans

Cluster: 0 mean: 1.0

Cluster: 1 mean: 1.9846

Cluster: 2 mean: 0.2097

Cluster: 3 mean: 1.6522

Cluster: 4 mean: 0.2824

Clustering Method: GMM

Cluster: 0 mean: 2.2667

Cluster: 1 mean: 0.1124

Cluster: 2 mean: 1.0385

Cluster: 3 mean: 1.093

Cluster: 4 mean: 0.2593

3. WSS

Hierarchical: 3151.18560967

KMeans: 2605.73353663

GMM: 2770.36699076

BSS

Hierarchical: 709.81439033

KMeans: 1255.26646337

GMM: 1090.63300924

4. One different attribute for each mode for each cluster.

See bellow for boxplots for each.

Hierarchical

Cluster 0 - Attr 5 - all same values (-0.411450), except one point (2.430427)

Cluster 1- Attr 5 - all same values (-0.411450), except one point (2.430427)

Cluster 2 - Attr 5 - same values, no distribution -0.411450

Cluster 3- Attr 5 - there's only one point here

Cluster 4 - Attr 5 - same values, no distribution 2.430427

KMeans

Cluster 0 - Attr 5 - same values, no distribution - 2.430427

Cluster 1- Attr 5 - all same values (-0.411450), except one point (2.430427)

Cluster 2 - Attr 5 - same values, no distribution - -0.41145

Cluster 3 - Attr 5 - same values, no distribution - -0.41145

Cluster 4 - Attr 5 - same values, no distribution - -0.41145

GMM

Cluster 0 - Attr 2- same values, no distribution - 0.87388

Cluster 1 - Attr 8 - all same values, no distribution - -0.696419

Cluster 2 - Attr 5 - same values, no distribution - -0.41145

Cluster 3 - Attr 5 - same values, no distribution - 2.430427

Cluster 4 - Attr 5 - same values, no distribution - -0.41145

```

In [13]: # Part 1 Cluster Count
from sklearn import mixture, cluster

clusterers = {
    'KMeans':cluster.KMeans(n_clusters=5, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=5, linkag
    'GMM':mixture.GMM(n_components=5, random_state=20160217),
    }

clusterings = {'True':heart_labels}

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

hiearchical_count = {}
kmeans_count = {}
gmm_count = {}

# Cluster sizes
for clusterer in clusterers.iterkeys():
    print "Clustering Method: " + clusterer
    for cluster in range(0, 5):
        bincount = np.bincount(clusterings[clusterer])
    for cluster in range(0, len(bincount)):
        print "Cluster: " + str(cluster) + " size: " + str(bincount[clu
        if (clusterer == 'KMeans'):
            kmeans_count[str(cluster)] = bincount[cluster]
        elif clusterer == 'Hierarchical':
            hiearchical_count[str(cluster)] = bincount[cluster]
        else:
            gmm_count[str(cluster)] = bincount[cluster]

```

```

Clustering Method: Hierarchical
Cluster: 0 size: 131
Cluster: 1 size: 129
Cluster: 2 size: 13
Cluster: 3 size: 1
Cluster: 4 size: 23
Clustering Method: KMeans
Cluster: 0 size: 39
Cluster: 1 size: 65
Cluster: 2 size: 62
Cluster: 3 size: 46
Cluster: 4 size: 85
Clustering Method: GMM
Cluster: 0 size: 60
Cluster: 1 size: 89
Cluster: 2 size: 78
Cluster: 3 size: 43
Cluster: 4 size: 27

```

```
In [14]: # Part 2 Cluster means
from sklearn import mixture, cluster

clusterers = {
    'KMeans':cluster.KMeans(n_clusters=5, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=5, linkag
    'GMM':mixture.GMM(n_components=5, random_state=20160217),
    }

clusterings = {'True':heart_labels}

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

hieararchical_means = {}
kmeans_means = {}
gmm_means = {}

# Cluster sizes
for clusterer in clusterers.iterkeys():
    print "Clustering Method: " + clusterer
    for cluster in range(0, 5):
        cluster_mean = 0
        cluster_count = 0
        for i in range(0, len(clusterings[clusterer])):
            if (clusterings[clusterer][i] == cluster):
                cluster_mean += heart_labels_values[i]
                cluster_count +=1

        cluster_mean = round(float(cluster_mean / float(cluster_count))
        if (clusterer == 'KMeans'):
            kmeans_means[str(cluster)] = cluster_mean
        elif clusterer == 'Hierarchical':
            hieararchical_means[str(cluster)] = cluster_mean
        else:
            gmm_means[str(cluster)] = cluster_mean

    print "Cluster: " + str(cluster) + " mean: " + str(cluster_mean
```

Clustering Method: Hierarchical

Cluster: 0 mean: 0.4275

Cluster: 1 mean: 1.1705

Cluster: 2 mean: 2.5385

Cluster: 3 mean: 0.0

Cluster: 4 mean: 1.7826

Clustering Method: KMeans

Cluster: 0 mean: 1.0

Cluster: 1 mean: 1.9846

Cluster: 2 mean: 0.2097

Cluster: 3 mean: 1.6522

Cluster: 4 mean: 0.2824

Clustering Method: GMM

Cluster: 0 mean: 2.2667

Cluster: 1 mean: 0.1124

Cluster: 2 mean: 1.0385

Cluster: 3 mean: 1.093

Cluster: 4 mean: 0.2593

In [15]:

```

# Part 3 WSS
from sklearn import mixture, cluster
pd.set_option('display.height', 1500)
pd.set_option('display.max_rows', 1500)

clusterers = {
    'KMeans':cluster.KMeans(n_clusters=5, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=5, linkag
    'GMM':mixture.GMM(n_components=5, random_state=20160217),
    }

clusterings = {'True':heart_labels}

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

def wss(kmeans):
    wss = 0
    # Iterate over each cluster
    for cluster in range(0, 5):
        cluster_instances = []
        # Iterate over each instance of the dataset
        for i in range(0, len(kmeans)):
            # If that dataset instance is in the right cluster
            if (kmeans[i] == cluster):
                cluster_instances.append(i)

        heart_cluster = heart_processed.iloc[cluster_instances]
        cols = heart_cluster.shape[1]

        for col in range(0, cols):
            feature_vals = heart_cluster[col].values
            attr_mean = 0
            for i in range(0, len(feature_vals)):
                attr_mean += feature_vals[i]

            attr_mean = float(attr_mean / float(len(feature_vals)))

            for i in range(0, len(feature_vals)):
                point = (feature_vals[i] - attr_mean)**2
                wss+= point
    return wss

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

for clusterer in clusterers.iterkeys():
    print str(clusterer) + " WSS:"
    if clusterer == 'KMeans':
        print wss(clusterings[clusterer])
    elif clusterer == 'Hierarchical':
        print wss(clusterings[clusterer])

```

```

else:
    print wss(clusterings[clusterer])

```

height has been deprecated.

```

Hierarchical WSS:
3151.18560967
KMeans WSS:
2605.73353663
GMM WSS:
2770.36699076

```

```

In [16]: # Part 3 BSS
from sklearn import mixture, cluster

def bss(kmeans):
    bss = 0
    mean_of_means = 0

    mean_of_means = heart_processed.mean()
    for cluster in range(0, 5):
        cluster_instances = []
        # Iterate over each instance of the dataset
        for i in range(0, len(kmeans)):
            # If that dataset instance is in the right cluster
            if (kmeans[i] == cluster):
                cluster_instances.append(i)

        heart_cluster = heart_processed.iloc[cluster_instances]
        attrs_means = heart_cluster.mean()

        for i in range(0, len(attrs_means)):
            bss += len(cluster_instances) * (attrs_means[i] - mean_of_m

    return bss

# print bss(hieararchical_means, hieararchical_count)
for clusterer in clusterers.iterkeys():
    print str(clusterer) + " BSS:"
    if clusterer == 'KMeans':
        print bss(clusterings[clusterer])
    elif clusterer == 'Hierarchical':
        print bss(clusterings[clusterer])
    else:
        print bss(clusterings[clusterer])

```

```

Hierarchical BSS:
709.81439033
KMeans BSS:
1255.26646337
GMM BSS:
1090.63300924

```

```
In [17]: ## Part #4 Attribute difference
from sklearn import mixture, cluster

clusterers = {
    'KMeans':cluster.KMeans(n_clusters=5, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=5, linkag
    'GMM':mixture.GMM(n_components=5, random_state=20160217),
    }

clusterings = {'True':heart_labels}

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

pnum = 1
for clusterer in clusterers.iterkeys():
    for i in range(0, 5):
        cluster_index = []
        for instance in range(0, len(clusterings[clusterer])):
            if i == clusterings[clusterer][instance]:
                cluster_index.append(instance)

        print str(clusterer) + " Cluster: " + str(i)

        df = pd.DataFrame(heart_processed.iloc[cluster_index])
        print df.head(10)

        fig = plt.figure()
        ax = plt.subplot(111)
        ax.boxplot(df.ix[:, i].values, positions = [i])


        df.boxplot(return_type='axes')

    #     ax.set_xlim(-0.5, 9.5)
    plt.show()
```

Hierarchical Cluster: 0

	0	1	2	3	4	5
6 \						
3	-1.941680	0.691095	-0.164289	-0.095506	0.051047	-0.411450
3419						
4	-1.498933	-1.446980	-1.202459	-0.095506	-0.835103	-0.411450
0199						
5	0.161372	0.691095	-1.202459	-0.659431	-0.218651	-0.411450
3419						
10	0.272059	0.691095	0.873880	0.468418	-1.066272	-0.411450
3419						
13	-1.166872	0.691095	-1.202459	-0.659431	0.301480	-0.411450
3419						
14	-0.281376	0.691095	-0.164289	2.272976	-0.931424	2.430427
3419						
15	0.272059	0.691095	-0.164289	1.032342	-1.528611	-0.411450
3419						
16	-0.724124	0.691095	-1.202459	-1.223355	-0.353500	-0.411450
3419						
17	-0.060002	0.691095	0.873880	0.468418	-0.160859	-0.411450
3419						
19	-0.613437	0.691095	-1.202459	-0.095506	0.359273	-0.411450
3419						


	7	8	9	10	11	12
3	1.633010	-0.696419	2.099753	2.264145	-0.721976	-0.894220
4	0.978071	-0.696419	0.295874	-0.976583	-0.721976	-0.894220
5	1.240047	-0.696419	-0.219520	-0.976583	-0.721976	-0.894220
10	-0.069831	-0.696419	-0.563116	0.643781	-0.721976	0.655877
13	1.021734	-0.696419	-0.906712	-0.976583	-0.721976	1.172577
14	0.541445	-0.696419	-0.477217	-0.976583	-0.721976	1.172577
15	1.065396	-0.696419	0.467672	-0.976583	-0.721976	-0.894220
16	0.803421	-0.696419	-0.047722	2.264145	-0.721976	1.172577
17	0.454120	-0.696419	0.124076	-0.976583	-0.721976	-0.894220
19	0.934409	-0.696419	-0.391318	-0.976583	-0.721976	-0.894220



Hierarchical Cluster: 1

	0	1	2	3	4	5
6 \						
1 0199	1.378929	0.691095	0.873880	1.596266	0.744555	-0.411450 1.01
2 0199	1.378929	0.691095	0.873880	-0.659431	-0.353500	-0.411450 1.01
7 3419	0.272059	-1.446980	0.873880	-0.659431	2.054515	-0.411450 -1.00
8 0199	0.936181	0.691095	0.873880	-0.095506	0.128103	-0.411450 1.01
11 0199	0.161372	-1.446980	-1.202459	0.468418	0.898668	-0.411450 1.01
18 3419	-0.724124	-1.446980	-0.164289	-0.095506	0.532650	-0.411450 -1.00
20 0199	1.046868	0.691095	-2.240629	-1.223355	-0.700254	-0.411450 1.01
21 0199	0.382746	-1.446980	-2.240629	1.032342	0.686763	2.430427 1.01
22 0199	0.382746	0.691095	-1.202459	-0.659431	0.706027	-0.411450 1.01
23 0199	0.382746	0.691095	-0.164289	0.017278	-0.449820	-0.411450 1.01


	7	8	9	10	11	12
1	-1.816334	1.435916	0.381773	0.643781	2.478425	-0.894220
2	-0.899420	1.435916	1.326662	0.643781	1.411625	1.172577
7	0.585108	1.435916	-0.391318	-0.976583	-0.721976	-0.894220
8	-0.113493	-0.696419	0.295874	0.643781	0.344824	1.172577
11	0.148482	-0.696419	0.209975	0.643781	-0.721976	-0.894220
18	-0.462794	-0.696419	-0.734914	-0.976583	-0.721976	-0.894220
20	-0.244481	1.435916	0.639470	0.643781	-0.721976	-0.894220
21	0.541445	-0.696419	-0.047722	-0.976583	-0.721976	-0.894220
22	0.454120	-0.696419	0.639470	0.643781	-0.721976	-0.894220
23	1.021734	-0.696419	1.842056	-0.976583	1.411625	1.172577



Hierarchical Cluster: 2

	0	1	2	3	4	5
6 \						
6	0.825494	-1.446980	0.873880	0.468418	0.397801	-0.41145
0199						
27	1.268242	-1.446980	-2.240629	1.032342	-0.411292	-0.41145
3419						
68	0.493433	0.691095	0.873880	2.160191	1.515120	-0.41145
0199						
90	0.825494	-1.446980	0.873880	1.596266	-1.605668	-0.41145
0199						
120	0.936181	-1.446980	0.873880	1.032342	3.075514	-0.41145
0199						
122	0.050685	0.691095	0.873880	0.468418	-0.584669	-0.41145
3419						
135	1.710989	0.691095	0.873880	0.750380	-1.413027	-0.41145
3419						
168	1.710989	0.691095	-0.164289	1.596266	0.417065	-0.41145
3419						
179	0.161372	-1.446980	0.873880	0.130063	3.114042	-0.41145
0199						
181	0.493433	0.691095	-2.240629	2.611330	0.436329	-0.41145
0199						


	7	8	9	10	11	12
6	0.454120	-0.696419	2.185652	2.264145	1.411625	-0.894220
27	-1.554358	-0.696419	1.326662	2.264145	-0.721976	-0.894220
68	-0.419131	1.435916	2.013854	2.264145	-0.721976	1.172577
90	-0.200818	-0.696419	4.419026	2.264145	2.478425	1.172577
120	0.192145	-0.696419	2.529248	0.643781	2.478425	1.172577
122	-1.685346	1.435916	3.903632	2.264145	-0.721976	1.172577
135	-1.074070	1.435916	1.326662	2.264145	-0.721976	1.172577
168	-1.641684	1.435916	1.584359	0.643781	0.344824	1.172577
179	0.017494	1.435916	0.725369	0.643781	1.411625	1.172577
181	-0.200818	-0.696419	2.701046	2.264145	-0.721976	1.172577



Hierarchical Cluster: 3

	0	1	2	3	4	5
6 \						
151	1.378929	-1.44698	-0.164289	-0.941393	6.099981	-0.41145
199						


	7	8	9	10	11	12
151	0.45412	-0.696419	0.467672	0.643781	-0.721976	1.172577



Hierarchical Cluster: 4

	0	1	2	3	4	5
6 \						
0	0.936181	0.691095	-2.240629	0.750380	-0.276443	2.430427
10199						
9	-0.170689	0.691095	0.873880	0.468418	-0.854367	2.430427
10199						
12	0.161372	0.691095	-0.164289	-0.095506	0.166631	2.430427
10199						
31	0.604120	0.691095	0.873880	-0.828608	-0.334236	2.430427
03419						
39	0.714807	0.691095	-0.164289	1.032342	-0.083802	2.430427
03419						
49	-0.170689	0.691095	-0.164289	-0.095506	-0.969952	2.430427
10199						
83	1.489615	0.691095	-0.164289	2.724115	0.513386	2.430427
10199						
110	0.161372	0.691095	0.873880	-0.377469	0.031783	2.430427
10199						
112	-1.277559	-1.446980	0.873880	0.017278	1.804082	2.430427
10199						
117	0.936181	0.691095	0.873880	-0.095506	1.592176	2.430427
10199						

	7	8	9	10	11	12
0	0.017494	-0.696419	1.068965	2.264145	-0.721976	0.655877
9	0.235807	1.435916	1.756157	2.264145	-0.721976	1.172577
12	-0.331806	1.435916	-0.391318	0.643781	0.344824	0.655877
31	0.454120	1.435916	0.295874	-0.976583	1.411625	1.172577
39	-0.550119	1.435916	-0.047722	0.643781	-0.721976	-0.894220
49	0.104820	-0.696419	0.124076	2.264145	-0.721976	-0.894220
83	0.017494	1.435916	0.467672	0.643781	-0.721976	1.172577
110	-0.244481	1.435916	0.124076	0.643781	0.344824	-0.894220
112	-0.593782	1.435916	1.670258	0.643781	-0.721976	1.172577
117	-0.768432	1.435916	0.639470	-0.976583	2.478425	1.172577




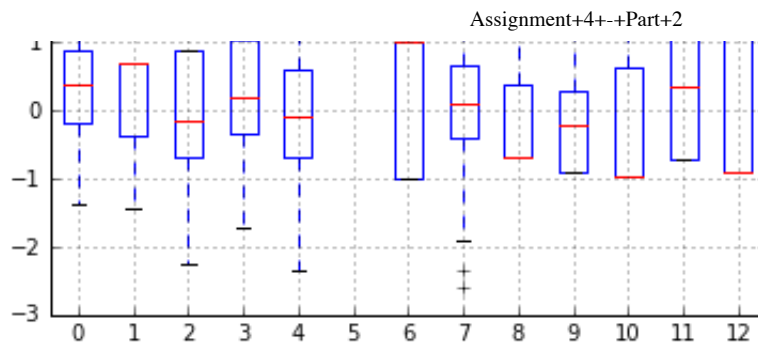


KMeans Cluster: 0

	0	1	2	3	4	5
6 \						
0	0.936181	0.691095	-2.240629	0.750380	-0.276443	2.430427
0199						
12	0.161372	0.691095	-0.164289	-0.095506	0.166631	2.430427
0199						
14	-0.281376	0.691095	-0.164289	2.272976	-0.931424	2.430427
3419						
21	0.382746	-1.446980	-2.240629	1.032342	0.686763	2.430427
0199						
31	0.604120	0.691095	0.873880	-0.828608	-0.334236	2.430427
3419						
39	0.714807	0.691095	-0.164289	1.032342	-0.083802	2.430427
3419						
43	0.493433	0.691095	-0.164289	1.032342	-0.680990	2.430427
3419						
48	1.157555	-1.446980	-0.164289	0.468418	3.268155	2.430427
0199						
49	-0.170689	0.691095	-0.164289	-0.095506	-0.969952	2.430427
0199						
63	-0.060002	-1.446980	-0.164289	0.186456	1.091309	2.430427
3419						


	7	8	9	10	11	12
0	0.017494	-0.696419	1.068965	2.264145	-0.721976	0.655877
12	-0.331806	1.435916	-0.391318	0.643781	0.344824	0.655877
14	0.541445	-0.696419	-0.477217	-0.976583	-0.721976	1.172577
21	0.541445	-0.696419	-0.047722	-0.976583	-0.721976	-0.894220
31	0.454120	1.435916	0.295874	-0.976583	1.411625	1.172577
39	-0.550119	1.435916	-0.047722	0.643781	-0.721976	-0.894220
43	0.323132	-0.696419	0.467672	-0.976583	-0.721976	-0.894220
48	0.323132	-0.696419	-0.219520	-0.976583	0.344824	-0.894220
49	0.104820	-0.696419	0.124076	2.264145	-0.721976	-0.894220
63	0.890746	-0.696419	-0.906712	-0.976583	-0.721976	-0.894220





KMeans Cluster: 1


	0	1	2	3	4	5	
6 \							
1	1.378929	0.691095	0.87388	1.596266	0.744555	-0.411450	1.010
199							
2	1.378929	0.691095	0.87388	-0.659431	-0.353500	-0.411450	1.010
199							
9	-0.170689	0.691095	0.87388	0.468418	-0.854367	2.430427	1.010
199							
24	0.604120	0.691095	0.87388	-0.095506	-0.796575	-0.411450	1.010
199							
29	-1.609620	0.691095	0.87388	-1.223355	-1.547875	-0.411450	1.010
199							
36	-1.277559	0.691095	0.87388	-0.659431	-1.355234	-0.411450	1.010
199							
37	0.272059	0.691095	0.87388	1.032342	0.551914	-0.411450	1.010
199							
38	0.050685	0.691095	0.87388	0.017278	2.035251	-0.411450	-1.003
419							
54	0.604120	0.691095	0.87388	-0.095506	0.108839	-0.411450	-1.003
419							
55	-0.060002	0.691095	0.87388	-0.433861	0.359273	-0.411450	1.010
199							
	7	8	9	10	11	12	
1	-1.816334	1.435916	0.381773	0.643781	2.478425	-0.894220	
2	-0.899420	1.435916	1.326662	0.643781	1.411625	1.172577	
9	0.235807	1.435916	1.756157	2.264145	-0.721976	1.172577	
24	-0.768432	1.435916	1.154864	0.643781	1.411625	1.172577	
29	-1.554358	1.435916	0.811268	0.643781	-0.721976	1.172577	
36	-1.292383	1.435916	1.240763	0.643781	-0.721976	1.172577	
37	-1.641684	1.435916	-0.391318	0.643781	0.344824	0.655877	
38	-0.768432	1.435916	0.124076	0.643781	0.344824	1.172577	
54	-0.244481	1.435916	0.295874	-0.976583	0.344824	1.172577	
55	-1.772671	1.435916	0.983066	0.643781	0.344824	1.172577	




KMeans Cluster: 2

	0	1	2	3	4	5
6 \						
7	0.272059	-1.446980	0.873880	-0.659431	2.054515	-0.41145
419						
11	0.161372	-1.446980	-1.202459	0.468418	0.898668	-0.41145
199						
18	-0.724124	-1.446980	-0.164289	-0.095506	0.532650	-0.41145
419						
25	-0.502750	-1.446980	-0.164289	-0.659431	-0.546141	-0.41145
419						
26	0.382746	-1.446980	-0.164289	-0.659431	1.784818	-0.41145
419						
27	1.268242	-1.446980	-2.240629	1.032342	-0.411292	-0.41145
419						
30	1.600302	-1.446980	-2.240629	0.468418	-0.160859	-0.41145
419						
32	1.046868	0.691095	-0.164289	0.468418	1.688497	-0.41145
419						
42	1.821676	-1.446980	-1.202459	1.596266	1.052781	-0.41145
419						
44	0.714807	-1.446980	0.873880	-0.095506	1.592176	-0.41145
199						


	7	8	9	10	11	12
7	0.585108	1.435916	-0.391318	-0.976583	-0.721976	-0.89422
11	0.148482	-0.696419	0.209975	0.643781	-0.721976	-0.89422
18	-0.462794	-0.696419	-0.734914	-0.976583	-0.721976	-0.89422
25	0.366795	-0.696419	0.467672	0.643781	-0.721976	-0.89422
26	0.978071	-0.696419	-0.906712	-0.976583	-0.721976	-0.89422
27	-1.554358	-0.696419	1.326662	2.264145	-0.721976	-0.89422
30	0.061157	-0.696419	0.639470	-0.976583	1.411625	-0.89422
32	0.366795	-0.696419	-0.906712	-0.976583	-0.721976	-0.89422
42	0.541445	-0.696419	-0.563116	-0.976583	1.411625	-0.89422
44	0.847083	-0.696419	-0.906712	-0.976583	-0.721976	-0.89422






KMeans Cluster: 3

	0	1	2	3	4	5	
6 \							
6	0.825494	-1.446980	0.873880	0.468418	0.397801	-0.41145	1.010
199							
8	0.936181	0.691095	0.873880	-0.095506	0.128103	-0.41145	1.010
199							
10	0.272059	0.691095	0.873880	0.468418	-1.066272	-0.41145	-1.003
419							
23	0.382746	0.691095	-0.164289	0.017278	-0.449820	-0.41145	1.010
199							
33	0.493433	0.691095	0.873880	0.186456	-0.257179	-0.41145	-1.003
419							
40	1.157555	-1.446980	0.873880	1.032342	-0.430556	-0.41145	1.010
199							
45	0.382746	0.691095	-0.164289	-1.110570	-0.334236	-0.41145	1.010
199							
47	-0.502750	0.691095	0.873880	1.032342	-0.083802	-0.41145	1.010
199							
51	1.157555	0.691095	0.873880	-0.659431	-1.355234	-0.41145	-1.003
419							
56	-0.502750	0.691095	-0.164289	0.468418	-0.276443	-0.41145	-1.003
419							
	7	8	9	10	11	12	
6	0.454120	-0.696419	2.185652	2.264145	1.411625	-0.894220	
8	-0.113493	-0.696419	0.295874	0.643781	0.344824	1.172577	
10	-0.069831	-0.696419	-0.563116	0.643781	-0.721976	0.655877	
23	1.021734	-0.696419	1.842056	-0.976583	1.411625	1.172577	
33	0.497783	-0.696419	-0.477217	0.643781	-0.721976	1.172577	
40	-1.554358	-0.696419	-0.047722	0.643781	2.478425	1.172577	
45	0.672433	-0.696419	1.240763	0.643781	0.344824	1.172577	
47	-0.943082	-0.696419	1.326662	0.643781	-0.721976	1.172577	
51	-0.419131	-0.696419	-0.563116	-0.976583	-0.721976	1.172577	
56	0.585108	-0.696419	-0.391318	0.643781	0.344824	1.172577	






KMeans Cluster: 4

	0	1	2	3	4	5
6 \						
3	-1.941680	0.691095	-0.164289	-0.095506	0.051047	-0.41145
419						
4	-1.498933	-1.446980	-1.202459	-0.095506	-0.835103	-0.41145
199						
5	0.161372	0.691095	-1.202459	-0.659431	-0.218651	-0.41145
419						
13	-1.166872	0.691095	-1.202459	-0.659431	0.301480	-0.41145
419						
15	0.272059	0.691095	-0.164289	1.032342	-1.528611	-0.41145
419						
16	-0.724124	0.691095	-1.202459	-1.223355	-0.353500	-0.41145
419						
17	-0.060002	0.691095	0.873880	0.468418	-0.160859	-0.41145
419						
19	-0.613437	0.691095	-1.202459	-0.095506	0.359273	-0.41145
419						
20	1.046868	0.691095	-2.240629	-1.223355	-0.700254	-0.41145
199						
22	0.382746	0.691095	-1.202459	-0.659431	0.706027	-0.41145
199						


	7	8	9	10	11	12
3	1.633010	-0.696419	2.099753	2.264145	-0.721976	-0.894220
4	0.978071	-0.696419	0.295874	-0.976583	-0.721976	-0.894220
5	1.240047	-0.696419	-0.219520	-0.976583	-0.721976	-0.894220
13	1.021734	-0.696419	-0.906712	-0.976583	-0.721976	1.172577
15	1.065396	-0.696419	0.467672	-0.976583	-0.721976	-0.894220
16	0.803421	-0.696419	-0.047722	2.264145	-0.721976	1.172577
17	0.454120	-0.696419	0.124076	-0.976583	-0.721976	-0.894220
19	0.934409	-0.696419	-0.391318	-0.976583	-0.721976	-0.894220
20	-0.244481	1.435916	0.639470	0.643781	-0.721976	-0.894220
22	0.454120	-0.696419	0.639470	0.643781	-0.721976	-0.894220




GMM Cluster: 0

	0	1	2	3	4	5
6 \						
1	1.378929	0.691095	0.87388	1.596266	0.744555	-0.41145
99						
2	1.378929	0.691095	0.87388	-0.659431	-0.353500	-0.41145
99						
6	0.825494	-1.446980	0.87388	0.468418	0.397801	-0.41145
99						
24	0.604120	0.691095	0.87388	-0.095506	-0.796575	-0.41145
99						
29	-1.609620	0.691095	0.87388	-1.223355	-1.547875	-0.41145
99						
36	-1.277559	0.691095	0.87388	-0.659431	-1.355234	-0.41145
99						
37	0.272059	0.691095	0.87388	1.032342	0.551914	-0.41145
99						
38	0.050685	0.691095	0.87388	0.017278	2.035251	-0.41145
19						
40	1.157555	-1.446980	0.87388	1.032342	-0.430556	-0.41145
99						
47	-0.502750	0.691095	0.87388	1.032342	-0.083802	-0.41145
99						

	7	8	9	10	11	12
1	-1.816334	1.435916	0.381773	0.643781	2.478425	-0.894220
2	-0.899420	1.435916	1.326662	0.643781	1.411625	1.172577
6	0.454120	-0.696419	2.185652	2.264145	1.411625	-0.894220
24	-0.768432	1.435916	1.154864	0.643781	1.411625	1.172577
29	-1.554358	1.435916	0.811268	0.643781	-0.721976	1.172577
36	-1.292383	1.435916	1.240763	0.643781	-0.721976	1.172577
37	-1.641684	1.435916	-0.391318	0.643781	0.344824	0.655877
38	-0.768432	1.435916	0.124076	0.643781	0.344824	1.172577
40	-1.554358	-0.696419	-0.047722	0.643781	2.478425	1.172577
47	-0.943082	-0.696419	1.326662	0.643781	-0.721976	1.172577






GMM Cluster: 1


	0	1	2	3	4	5
6 \						
3	-1.941680	0.691095	-0.164289	-0.095506	0.051047	-0.41145
419						
4	-1.498933	-1.446980	-1.202459	-0.095506	-0.835103	-0.41145
199						
5	0.161372	0.691095	-1.202459	-0.659431	-0.218651	-0.41145
419						
11	0.161372	-1.446980	-1.202459	0.468418	0.898668	-0.41145
199						
15	0.272059	0.691095	-0.164289	1.032342	-1.528611	-0.41145
419						
17	-0.060002	0.691095	0.873880	0.468418	-0.160859	-0.41145
419						
18	-0.724124	-1.446980	-0.164289	-0.095506	0.532650	-0.41145
419						
19	-0.613437	0.691095	-1.202459	-0.095506	0.359273	-0.41145
419						
22	0.382746	0.691095	-1.202459	-0.659431	0.706027	-0.41145
199						
25	-0.502750	-1.446980	-0.164289	-0.659431	-0.546141	-0.41145
419						

	7	8	9	10	11	12
3	1.633010	-0.696419	2.099753	2.264145	-0.721976	-0.89422
4	0.978071	-0.696419	0.295874	-0.976583	-0.721976	-0.89422
5	1.240047	-0.696419	-0.219520	-0.976583	-0.721976	-0.89422
11	0.148482	-0.696419	0.209975	0.643781	-0.721976	-0.89422
15	1.065396	-0.696419	0.467672	-0.976583	-0.721976	-0.89422
17	0.454120	-0.696419	0.124076	-0.976583	-0.721976	-0.89422
18	-0.462794	-0.696419	-0.734914	-0.976583	-0.721976	-0.89422
19	0.934409	-0.696419	-0.391318	-0.976583	-0.721976	-0.89422
22	0.454120	-0.696419	0.639470	0.643781	-0.721976	-0.89422
25	0.366795	-0.696419	0.467672	0.643781	-0.721976	-0.89422



GMM Cluster: 2

	0	1	2	3	4	5
6 \						
8	0.936181	0.691095	0.873880	-0.095506	0.128103	-0.41145
199						
10	0.272059	0.691095	0.873880	0.468418	-1.066272	-0.41145
419						
13	-1.166872	0.691095	-1.202459	-0.659431	0.301480	-0.41145
419						
16	-0.724124	0.691095	-1.202459	-1.223355	-0.353500	-0.41145
419						
20	1.046868	0.691095	-2.240629	-1.223355	-0.700254	-0.41145
199						
23	0.382746	0.691095	-0.164289	0.017278	-0.449820	-0.41145
199						
33	0.493433	0.691095	0.873880	0.186456	-0.257179	-0.41145
419						
34	-1.166872	0.691095	-0.164289	-0.095506	-0.276443	-0.41145
419						
41	-1.609620	0.691095	-2.240629	0.468418	-0.931424	-0.41145
419						
45	0.382746	0.691095	-0.164289	-1.110570	-0.334236	-0.41145
199						
	7	8	9	10	11	12
8	-0.113493	-0.696419	0.295874	0.643781	0.344824	1.172577
10	-0.069831	-0.696419	-0.563116	0.643781	-0.721976	0.655877
13	1.021734	-0.696419	-0.906712	-0.976583	-0.721976	1.172577
16	0.803421	-0.696419	-0.047722	2.264145	-0.721976	1.172577
20	-0.244481	1.435916	0.639470	0.643781	-0.721976	-0.894220
23	1.021734	-0.696419	1.842056	-0.976583	1.411625	1.172577
33	0.497783	-0.696419	-0.477217	0.643781	-0.721976	1.172577
34	1.283709	1.435916	-0.563116	-0.976583	-0.721976	-0.894220
41	1.240047	1.435916	0.295874	-0.976583	-0.721976	1.172577
45	0.672433	-0.696419	1.240763	0.643781	0.344824	1.172577




GMM Cluster: 3

	0	1	2	3	4	5
6 \						
0	0.936181	0.691095	-2.240629	0.750380	-0.276443	2.430427
0199						
9	-0.170689	0.691095	0.873880	0.468418	-0.854367	2.430427
0199						
12	0.161372	0.691095	-0.164289	-0.095506	0.166631	2.430427
0199						
14	-0.281376	0.691095	-0.164289	2.272976	-0.931424	2.430427
3419						
21	0.382746	-1.446980	-2.240629	1.032342	0.686763	2.430427
0199						
31	0.604120	0.691095	0.873880	-0.828608	-0.334236	2.430427
3419						
39	0.714807	0.691095	-0.164289	1.032342	-0.083802	2.430427
3419						
43	0.493433	0.691095	-0.164289	1.032342	-0.680990	2.430427
3419						
48	1.157555	-1.446980	-0.164289	0.468418	3.268155	2.430427
0199						

49 -0.170689 0.691095 -0.164289 -0.095506 -0.969952 2.430427 1.010199

	7	8	9	10	11	12
0	0.017494	-0.696419	1.068965	2.264145	-0.721976	0.655877
9	0.235807	1.435916	1.756157	2.264145	-0.721976	1.172577
12	-0.331806	1.435916	-0.391318	0.643781	0.344824	0.655877
14	0.541445	-0.696419	-0.477217	-0.976583	-0.721976	1.172577
21	0.541445	-0.696419	-0.047722	-0.976583	-0.721976	-0.894220
31	0.454120	1.435916	0.295874	-0.976583	1.411625	1.172577
39	-0.550119	1.435916	-0.047722	0.643781	-0.721976	-0.894220
43	0.323132	-0.696419	0.467672	-0.976583	-0.721976	-0.894220
48	0.323132	-0.696419	-0.219520	-0.976583	0.344824	-0.894220
49	0.104820	-0.696419	0.124076	2.264145	-0.721976	-0.894220




GMM Cluster: 4

	0	1	2	3	4	5
6 \						
7	0.272059	-1.44698	0.873880	-0.659431	2.054515	-0.41145
419						
30	1.600302	-1.44698	-2.240629	0.468418	-0.160859	-0.41145
419						
42	1.821676	-1.44698	-1.202459	1.596266	1.052781	-0.41145
419						
50	-1.498933	-1.44698	-1.202459	-1.505317	-0.950688	-0.41145
419						
61	-0.945498	-1.44698	-0.164289	0.581203	-1.355234	-0.41145
199						
77	-0.392063	-1.44698	-0.164289	0.468418	1.168366	-0.41145
199						
101	0.272059	-1.44698	0.873880	-0.208291	1.072045	-0.41145
199						
113	0.825494	-1.44698	-0.164289	-0.095506	0.301480	-0.41145
419						

```
148 0.604120 -1.44698 -0.164289 -1.674494 1.361007 -0.41145 -1.003
419
151 1.378929 -1.44698 -0.164289 -0.941393 6.099981 -0.41145 1.010
199
```

	7	8	9	10	11	12
7	0.585108	1.435916	-0.391318	-0.976583	-0.721976	-0.894220
30	0.061157	-0.696419	0.639470	-0.976583	1.411625	-0.894220
42	0.541445	-0.696419	-0.563116	-0.976583	1.411625	-0.894220
50	0.803421	-0.696419	-0.906712	-0.976583	0.344824	-0.894220
61	0.454120	1.435916	0.295874	2.264145	-0.721976	-0.894220
77	-0.331806	-0.696419	0.381773	-0.976583	0.344824	-0.894220
101	0.410458	-0.696419	-0.906712	-0.976583	0.344824	-0.894220
113	-2.296622	-0.696419	0.124076	0.643781	0.344824	1.172577
148	0.454120	-0.696419	-0.906712	-0.976583	0.344824	-0.894220
151	0.454120	-0.696419	0.467672	0.643781	-0.721976	1.172577



In [18]:

```

## Part #5 Plotting boxplots of attributes
from sklearn import mixture, cluster

clusterers = {
    'KMeans':cluster.KMeans(n_clusters=5, init='random', random_state=2
    'Hierarchical':cluster.AgglomerativeClustering(n_clusters=5, linkag
    'GMM':mixture.GMM(n_components=5, random_state=20160217),
    }

clusterings = {'True':heart_labels}

for clusterer in clusterers.iterkeys():
    clusterings[clusterer]=clusterers[clusterer].fit_predict(heart_proc

all_values = []

for clusterer in clusterers.iterkeys():
    print "Model: " + clusterer
    for i in range(0, 5):
        cluster_index = []
        for instance in range(0, len(clusterings[clusterer])):
            if i == clusterings[clusterer][instance]:
                cluster_index.append(instance)

        df = pd.DataFrame(heart_processed.iloc[cluster_index])
        #         print "Cluster: " + str(i)
        #         print df.head(10)
        all_values.append(df)

    all_values.append(pd.DataFrame(heart_processed.iloc[clusterings[clu

font = 20
small_font = 20
f, ax = plt.subplots(3, 6, figsize=(25,15))
ax[0][0].boxplot(all_values[0].values)
ax[0][0].set_title("Cluster 0", fontsize=small_font)
ax[0][0].set_ylabel("Hierarchical Clustering", fontsize=font)

ax[0][1].boxplot(all_values[1].values)
ax[0][1].set_title("Cluster 1", fontsize=small_font)

ax[0][2].boxplot(all_values[2].values)
ax[0][2].set_title("Cluster 2", fontsize=small_font)

ax[0][3].boxplot(all_values[3].values)
ax[0][3].set_title("Cluster 3", fontsize=small_font)

ax[0][4].boxplot(all_values[4].values)
ax[0][4].set_title("Cluster 4", fontsize=small_font)

ax[0][5].boxplot(all_values[5].values)
ax[0][5].set_title("All Data", fontsize=small_font)

ax[1][0].boxplot(all_values[6].values)

```

```

ax[1][0].set_ylabel("KMeans", fontsize=font)
ax[1][1].boxplot(all_values[7].values)
ax[1][2].boxplot(all_values[8].values)
ax[1][3].boxplot(all_values[9].values)
ax[1][4].boxplot(all_values[10].values)
ax[1][5].boxplot(all_values[11].values)

ax[2][0].boxplot(all_values[11].values)
ax[2][0].set_ylabel("GMM", fontsize=font)
ax[2][1].boxplot(all_values[12].values)
ax[2][2].boxplot(all_values[13].values)
ax[2][3].boxplot(all_values[14].values)
ax[2][4].boxplot(all_values[15].values)
ax[2][5].boxplot(all_values[16].values)

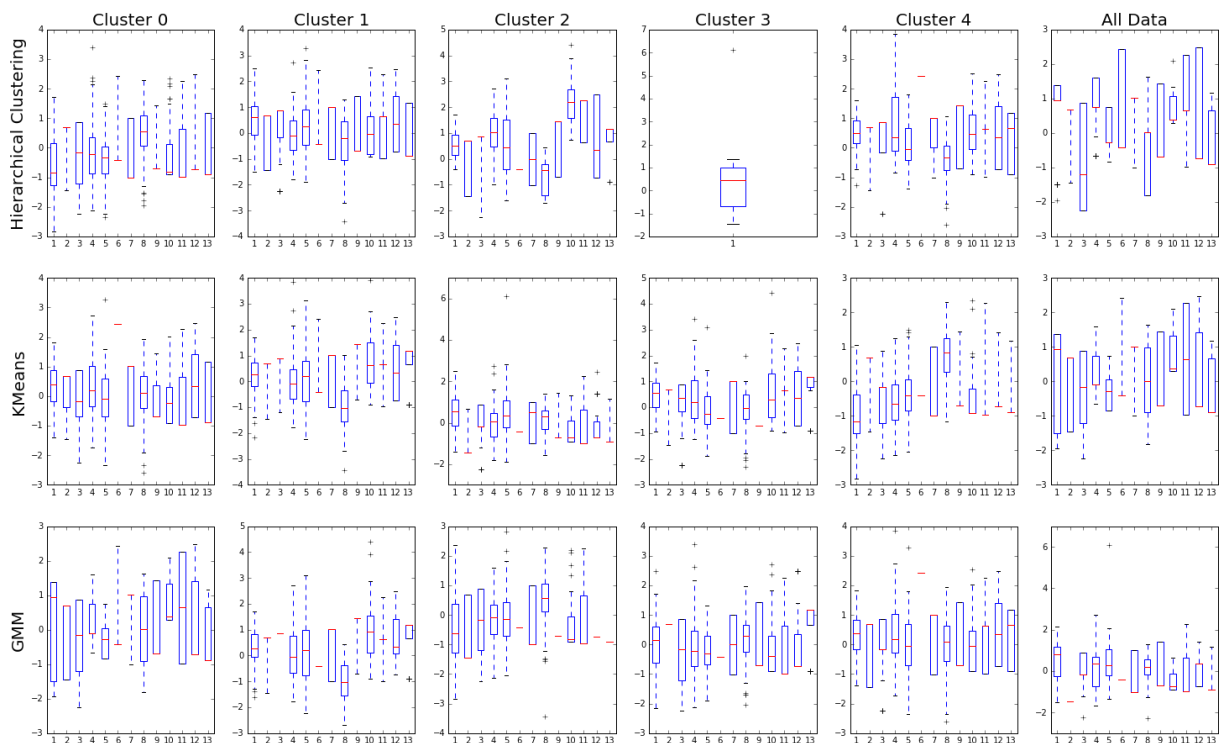
plt.show()

```

Model: Hierarchical

Model: KMeans

Model: GMM



Question 3: Clustering Census Data (15 points)

This question is a bit trickier. You'll be clustering data with a mix of numeric and nominal attributes. I've encoded the one ordinal attribute for you. You'll be looking at a few different techniques to work with this sort of data.

1. Using `census_data_scaled`, run KMeans and GMM clustering, while increasing the number of clusters to 100 taking steps of 4 (e.g. `range(1,100,4)`). Plot the WSS with

- respect to the number of clusters. It may also be helpful to plot the change in WSS as the number of clusters increase. Where does the WSS error seem to plateau?
2. One of the issues we've noted is that distance functions often behave nonintuitively in very high dimensions. Let's try to address this problem in a couple of ways. One option is to project into a lower dimensional space. We used PCA for this in Assignment 1.
 - Project the census data into a 3-dimensional space and a 7-dimensional space using PCA.
 - Repeat the clustering experiment, plotting the WSS for both settings. Do you notice a difference in how the WSS decreases?
 - Choose a value for the number of clusters based on the plot. Pick one of the clusters from the 7D decomposition in PCA space and examine that cluster in the original attribute space. What are the salient attribute values for that cluster?
 3. In some cases, the data is naturally high-dimensional. However, in this situation the high dimensionality of the data is the consequence of our preprocessing. Many of the attributes in the data arise from nominal attributes that have been translated into a OneHot encoding. If these nominal attributes could be encoded more densely, our clusterings might be more useful. We can try to encode the nominal attributes with many values (WorkClass, MaritalStatus, Occupation, NativeCountry) in the original data differently. You could, for example, encode Marital Status as Married or Single, re-encode countries into continents, re-encode WorkClass into Government/non-Government, and Occupation into white-collar or blue-collar. These are just examples, and you can choose your own encoding strategy.
 - Re-encode the nominal attributes using a scheme of your choice
 - Repeat the experiment, plotting the WSS again. What do you notice?
 - Choose a number of clusters based on the WSS plot. Pick one cluster and report the salient attributes for that cluster.

Answers 3

1. The WSS error seems to plateau around 70 clusters for GMM and around 80 for KMeans
2. The WSS values decrease very differently than the high dimensionality data. Furthermore, between the 3 and 7 D data, the error decreases at the same speed, but at different values. The shape is the same while the actual error values are different. Higher for the 7D data.
Also when performing PCA the error starts to plateau much earlier, at around 10 clusters. Furthermore, the lower the dimensionality of the data, the lower the WSS error.
For PCA 3 dimensions, the error starts out at 300,000 and goes down to 9000 for each
For PCA 7 dimensions, the error starts out at around 600,000 and goes down to about 50,000, so not as low as the lower dimensional PCA 3.

Plotting one of the clusters of the 7D space and the equivalent cluster with 88 features, we

see that only few features have a good feature value distribution. Most features in the original dataset are have a very small range of values. There are 9 features which are relevant: 0, 3, 4, 5, 6, 54, 74, 77 and 82. We see that the choice of a 7D PCA is good, though we could have probably increased it to 9 capture the 2 additional features that might have contributed to the clustering.

3. Reencoding the high dimensional nominal values and plotting the WSS we notice that the plot starts to look very similar to the plot of the PCA WSS error. This makes sense, because with the reenconding we are doing the same as what PCA does, reduces the dimensionality of the attributes. Thus the simiar shape makes sense.

It's important to note that the WSS error is much higher than the PCA, yet lower than performing one-hot encoding.

Salient attributes: 0, 3, 4, 8, 10, 57, 70, 72, 73, 79, 80, 83

```

In [19]: census_data = urllib2.urlopen("http://archive.ics.uci.edu/ml/machine-le
census_orig = pd.read_csv(census_data, quotechar='"', skipinitialspace=
                        names=['Age', 'WorkClass', 'FnlWgt', 'Edu
                        'Occupation', 'Relationship', 'Ra
                        'CapitalGain', 'CapitalLoss', 'Ho
                        'NativeCountry', 'Label'],
                        na_values="?", index_col=False)

census_orig = census_orig.dropna()
census_orig = census_orig.drop('FnlWgt', 1)
education_translation = {'Preschool':0,
                        '1st-4th':1,
                        '5th-6th':2,
                        '7th-8th':3,
                        '9th':4,
                        '10th':5,
                        '11th':6,
                        '12th':7,
                        'HS-grad':8,
                        'Some-college':9,
                        'Assoc-voc':10,
                        'Assoc-acdm':11,
                        'Bachelors':12,
                        'Masters':13,
                        'Prof-school':14,
                        'Doctorate':15
                        }
census_orig['Education'] = census_orig['Education'].apply(lambda x: edu

# Convert labels from strings to boolean
label_encoder = preprocessing.LabelEncoder()
census_labels = pd.DataFrame(label_encoder.fit_transform(census_orig.il

# Convert nominal attributes to encoded versions
attr_encoder = feature_extraction.DictVectorizer(sparse=False)
census_data = pd.DataFrame(attr_encoder.fit_transform(census_orig.iloc[
census_data.columns = attr_encoder.get_feature_names()

census_data_scaled = pd.DataFrame(preprocessing.StandardScaler().fit_tr
#census_data_scaled.columns = attr_encoder.get_feature_names()

```

```

In [20]: # Part 1
from sklearn import mixture, cluster

def wss(kmeans, cluster_size):
    wss = 0
    # Iterate over each cluster
    for cluster in range(0, cluster_size):
        cluster_instances = []
        # Iterate over each instance of the dataset
        for i in range(0, len(kmeans)):
            # If that dataset instance is in the right cluster
            if (kmeans[i] == cluster):
                cluster_instances.append(i)

        census_cluster = census_data_scaled.iloc[cluster_instances]
        cols = census_cluster.shape[1]

        for col in range(0, cols):
            feature_vals = census_cluster[col].values
            attr_mean = 0
            for i in range(0, len(feature_vals)):
                attr_mean += feature_vals[i]

            attr_mean = float(attr_mean / float(len(feature_vals)))

            for i in range(0, len(feature_vals)):
                point = (feature_vals[i] - attr_mean)**2
                wss+= point

    return wss

cluster_size = []
kmeans_wss_vals = []
gmm_wss_vals = []
for cluster_num in range(1, 100, 4):
    km = cluster.KMeans(n_clusters=cluster_num, init='random', random_s
    gmm = mixture.GMM(n_components=cluster_num, random_state=20160217)

    kmeans_wss = wss(km.fit_predict(census_data_scaled), cluster_num)
    gmm_wss = wss(gmm.fit_predict(census_data_scaled), cluster_num)

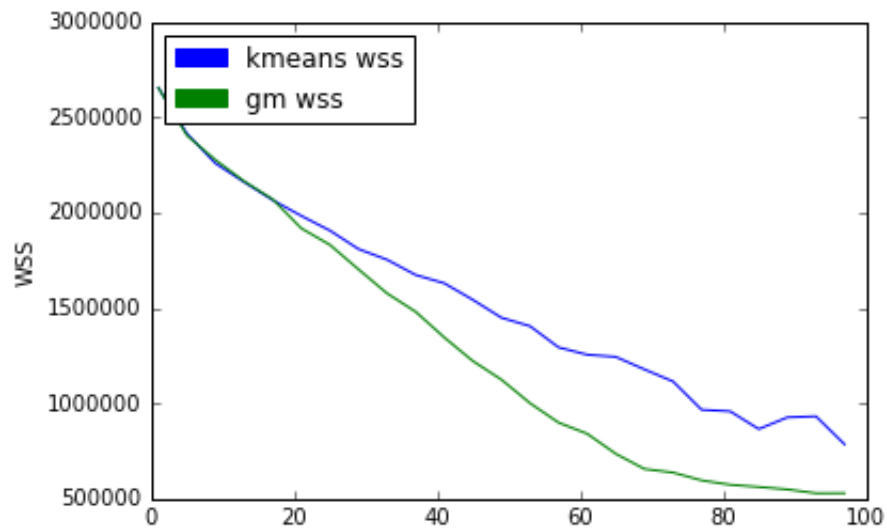
    cluster_size.append(cluster_num)
    kmeans_wss_vals.append(kmeans_wss)
    gmm_wss_vals.append(gmm_wss)

print "Cluster: " + str(cluster_num)
print "KM wss : " + str(kmeans_wss) + " GMM wss: " + str(gmm_wss)

```

Cluster: 1
KM wss : 2654256.00002 GMM wss: 2654256.00002
Cluster: 5
KM wss : 2413587.05166 GMM wss: 2403845.90346
Cluster: 9
KM wss : 2257288.92756 GMM wss: 2275604.63554
Cluster: 13
KM wss : 2159732.92708 GMM wss: 2164168.37649
Cluster: 17
KM wss : 2065185.73405 GMM wss: 2069432.46524
Cluster: 21
KM wss : 1982692.4696 GMM wss: 1917221.74872
Cluster: 25
KM wss : 1905639.43897 GMM wss: 1830700.96148
Cluster: 29
KM wss : 1808786.1577 GMM wss: 1702252.06306
Cluster: 33
KM wss : 1752780.28775 GMM wss: 1576220.37798
Cluster: 37
KM wss : 1672867.74956 GMM wss: 1479942.97078
Cluster: 41
KM wss : 1629500.98015 GMM wss: 1345238.04516
Cluster: 45
KM wss : 1542574.00487 GMM wss: 1221620.49957
Cluster: 49
KM wss : 1448685.99897 GMM wss: 1123931.39958
Cluster: 53
KM wss : 1404466.77561 GMM wss: 1000544.07076
Cluster: 57
KM wss : 1293064.34984 GMM wss: 898638.329022
Cluster: 61
KM wss : 1254866.25114 GMM wss: 840297.592565
Cluster: 65
KM wss : 1243397.80166 GMM wss: 735228.388188
Cluster: 69
KM wss : 1177783.52734 GMM wss: 655924.623116
Cluster: 73
KM wss : 1114758.41641 GMM wss: 637116.979107
Cluster: 77
KM wss : 966211.433307 GMM wss: 595922.428696
Cluster: 81
KM wss : 958060.178007 GMM wss: 573118.084469
Cluster: 85
KM wss : 864980.122724 GMM wss: 561486.245808
Cluster: 89
KM wss : 926380.048568 GMM wss: 548736.064057
Cluster: 93
KM wss : 931768.420137 GMM wss: 528466.64443
Cluster: 97
KM wss : 784998.311987 GMM wss: 528959.304985

```
In [21]: plt.plot(cluster_size, kmeans_wss_vals, 'b', cluster_size, gmm_wss_vals  
blue_patch = mpatches.Patch(color='blue', label='kmeans wss')  
green_patch = mpatches.Patch(color='green', label='gm wss')  
  
plt.legend(handles=[blue_patch, green_patch], loc='upper left')  
  
plt.ylabel("WSS")  
plt.show()
```



In [22]:

```

# Part 2
from sklearn.decomposition import PCA

def wss(kmeans, cluster_size, census_data_scaled):
    wss = 0
    # Iterate over each cluster
    for cluster in range(0, cluster_size):
        cluster_instances = []
        # Iterate over each instance of the dataset
        for i in range(0, len(kmeans)):
            # If that dataset instance is in the right cluster
            if (kmeans[i] == cluster):
                cluster_instances.append(i)

        census_cluster = census_data_scaled.iloc[cluster_instances]
        cols = census_cluster.shape[1]

        for col in range(0, cols):
            feature_vals = census_cluster[col].values
            attr_mean = 0
            for i in range(0, len(feature_vals)):
                attr_mean += feature_vals[i]

            if len(feature_vals) == 0:
                attr_mean = 0
            else:
                attr_mean = float(attr_mean / float(len(feature_vals)))

            for i in range(0, len(feature_vals)):
                point = (feature_vals[i] - attr_mean)**2
                wss += point

    return wss

pca_3 = pd.DataFrame(PCA(n_components=3).fit_transform(census_data_scaled))
pca_7 = pd.DataFrame(PCA(n_components=7).fit_transform(census_data_scaled))

pca_list = [pca_3, pca_7]

for pca_data in pca_list:

    pca_cluster_size = []
    pca_kmeans_wss_vals = []
    pca_gmm_wss_vals = []
    print pca_data.head()
    print "-----PCA VALUES-----"
    for cluster_num in range(1, 100, 4):
        km = cluster.KMeans(n_clusters=cluster_num, init='random', random_state=201602)
        gmm = mixture.GMM(n_components=cluster_num, random_state=201602)

        kmeans_wss = wss(km.fit_predict(pca_data), cluster_num, pca_data)
        gmm_wss = wss(gmm.fit_predict(pca_data), cluster_num, pca_data)

```



```
pca_cluster_size.append(cluster_num)
pca_kmeans_wss_vals.append(kmeans_wss)
pca_gmm_wss_vals.append(gmm_wss)

print "Cluster: " + str(cluster_num)
print "KM wss : " + str(kmeans_wss) + " GMM wss: " + str(gmm_ws

plt.plot(pca_cluster_size, pca_kmeans_wss_vals, 'b', pca_cluster_si
blue_patch = mpatches.Patch(color='blue', label='kmeans wss')
green_patch = mpatches.Patch(color='green', label='gm wss')

plt.legend(handles=[blue_patch, green_patch], loc='upper left')

plt.ylabel("run time in seconds")
plt.show()
```

	0	1	2
0	0.607939	-2.269147	-0.763535
1	-2.788245	-1.317830	-0.023245
2	0.803568	0.708483	-1.199120
3	-0.901320	2.529662	1.173713
4	1.204012	-1.421205	4.418078

-----PCA VALUES-----

Cluster: 1
KM wss : 304659.136888 GMM wss: 304659.136888

Cluster: 5
KM wss : 96032.5125022 GMM wss: 123704.17403

Cluster: 9
KM wss : 57337.0908372 GMM wss: 77060.8308305

Cluster: 13
KM wss : 42569.6730684 GMM wss: 59749.8019737

Cluster: 17
KM wss : 34404.8244546 GMM wss: 45424.5730271

Cluster: 21
KM wss : 28910.7311625 GMM wss: 39878.0591016

Cluster: 25
KM wss : 25178.0482025 GMM wss: 37192.1741238

Cluster: 29
KM wss : 22388.6911411 GMM wss: 33357.8453997

Cluster: 33
KM wss : 20720.9293346 GMM wss: 31064.6806697

Cluster: 37
KM wss : 18869.0939049 GMM wss: 29592.1449597

Cluster: 41
KM wss : 17442.4623923 GMM wss: 27061.9355364

Cluster: 45
KM wss : 16289.070252 GMM wss: 26403.8788113

Cluster: 49
KM wss : 15342.8269028 GMM wss: 25622.0442064

Cluster: 53
KM wss : 14523.5707016 GMM wss: 23377.8640907

Cluster: 57
KM wss : 13770.5697844 GMM wss: 23164.7934616

Cluster: 61
KM wss : 13126.232313 GMM wss: 22890.0196648

Cluster: 65
KM wss : 12652.3087242 GMM wss: 19910.7798498

Cluster: 69
KM wss : 12169.8127832 GMM wss: 20461.5181309

Cluster: 73
KM wss : 11735.8321891 GMM wss: 20320.4479115

Cluster: 77
KM wss : 11417.8883367 GMM wss: 18680.9543263

Cluster: 81
KM wss : 11030.1787208 GMM wss: 18109.7412403

Cluster: 85
KM wss : 10717.4728204 GMM wss: 16978.6376745

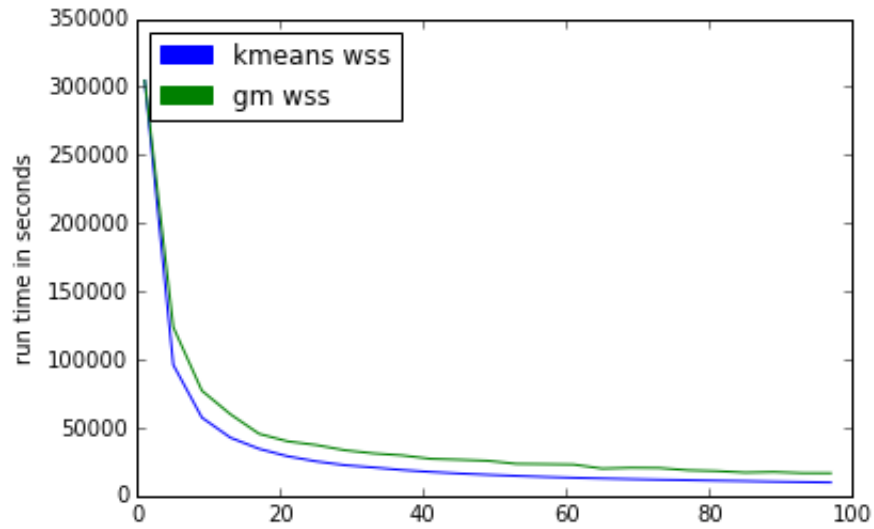
Cluster: 89
KM wss : 10312.2236635 GMM wss: 17390.6841564

Cluster: 93
KM wss : 10000.1615500 GMM wss: 16601.0047070

KM wss : 10000.1615502 GMM wss: 16621.9947978

Cluster: 97

KM wss : 9737.05356624 GMM wss: 16575.505536



	0	1	2	3	4	5
6						
0	0.607939	-2.269147	-0.763535	1.360590	-0.782690	-1.889694
049						
1	-2.788245	-1.317830	-0.023245	-0.209085	-0.328010	-0.509259
226						
2	0.803568	0.708483	-1.199120	-0.276066	0.336155	-1.341190
172						
3	-0.901320	2.529662	1.173713	-0.648880	-3.038289	1.467002
316						
4	1.204012	-1.421205	4.418078	0.123548	1.472908	2.488223
141						

-----PCA VALUES-----

Cluster: 1

KM wss : 531780.675644 GMM wss: 531780.675644

Cluster: 5

KM wss : 288863.707382 GMM wss: 332772.115569

Cluster: 9

KM wss : 194777.021722 GMM wss: 214440.735736

Cluster: 13

KM wss : 156218.276029 GMM wss: 181784.370583

Cluster: 17


KM wss : 133779.584715 GMM wss: 158381.643779


Cluster: 21

KM wss : 118918.354128 GMM wss: 142873.48907

Cluster: 25

KM wss : 108698.145359 GMM wss: 128234.688623





```

In [23]: # Part 2 salient attribute
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

pca_data = pd.DataFrame(PCA(n_components=7).fit_transform(census_data_s
cluster_size = 40

gmm = mixture.GMM(n_components=cluster_size, random_state=20160217)

gmm_predict_pca = gmm.fit_predict(pca_data)

cluster_instances = []
# Iterate over each instance of the dataset
for i in range(0, len(gmm_predict_pca)):
    # If that dataset instance is in the right cluster
    if (gmm_predict_pca[i] == 5):
        cluster_instances.append(i)


pca_cluster = pca_data.iloc[cluster_instances]
normal_cluster = census_data_scaled.iloc[cluster_instances]


plt.figure(figsize=(20,20))

normal_cluster.boxplot(return_type='axes')

```


Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x112891e10>





```
In [24]: # Part 2 PCA boxplot
pca_cluster.boxplot(return_type='axes')
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x111eafb0>
```



```
In [25]: # Part 3
census_data = urllib2.urlopen("http://archive.ics.uci.edu/ml/machine-le
census_orig = pd.read_csv(census_data, quotechar='"', skipinitialspace=
                        names=['Age', 'WorkClass', 'FnlWgt', 'Edu
                        'Occupation', 'Relationship', 'Ra
                        'CapitalGain', 'CapitalLoss', 'Ho
                        'NativeCountry', 'Label'],
                        na_values="?", index_col=False)

census_orig = census_orig.dropna()
census_orig = census_orig.drop('FnlWgt', 1)
education_translation = {'Preschool': 0,
                        '1st-4th': 1,
                        '5th-6th': 2,
                        '7th-8th': 3,
                        '9th': 4,
                        '10th': 5,
                        '11th': 6,
                        '12th': 7,
                        'HS-grad': 8,
                        'Some-college': 9,
                        'Assoc-voc': 10,
                        'Assoc-acdm': 11,
                        'Bachelors': 12,
                        'Masters': 13,
                        'Prof-school': 14,
                        'Doctorate': 15
                        }
census_orig['Education'] = census_orig['Education'].apply(lambda x: edu
```


In [26]:

```

marital_status_translation = {'Separated':0,
                              'Widowed':0,
                              'Divorced':0,
                              'Married-spouse-absent':1,
                              'Never-married':0,
                              'Married-AF-spouse':1,
                              'Married-civ-spouse':1
                              }

relationship_translation = {'Own-child':1,
                            'Wife':1,
                            'Unmarried':0,
                            'Other-relative':1,
                            'Husband':1,
                            'Not-in-family': 0
                            }

work_class_translation = {'Self-emp-inc':'Non-Government',
                          'State-gov':'Government',
                          'Without-pay':'Non-Government',
                          'Private':'Non-Government',
                          'Local-gov':'Government',
                          'Self-emp-not-inc':'Non-Government',
                          'Federal-gov':'Government'
                          }

occupation_translation = {'Farming-fishing':'blue-collar',
                          'Armed-Forces':'blue-collar',
                          'Craft-repair':'blue-collar',
                          'Other-service':'white-collar',
                          'Transport-moving':'blue-collar',
                          'Prof-specialty':'white-collar',
                          'Sales':'white-collar',
                          'Exec-managerial':'white-collar',
                          'Handlers-cleaners':'blue-collar',
                          'Adm-clerical':'white-collar',
                          'Protective-serv':'white-collar',
                          'Tech-support':'white-collar',
                          'Priv-house-serv':'blue-collar',
                          'Machine-op-inspct':'white-collar'
                          }

native_country_transformation = {'Canada': 'NA',
                                  'Hong': 'ASIA',
                                  'Dominican-Republic': 'CA',
                                  'Italy': 'EUROPE',
                                  'Ireland': 'EUROPE',
                                  'Outlying-US(Guam-USVI-etc)': 'ASIA',
                                  'Scotland': 'EUROPE',
                                  'Cambodia': 'ASIA',
                                  'France': 'EUROPE',
                                  'Peru': 'SA',
                                  'Laos': 'ASIA',
                                  'Ecuador': 'CA',
                                  'Iran': 'ASIA'.

```

```

'Cuba': 'CA',
'Guatemala': 'SA',
'Germany': 'EUROPE',
'Thailand': 'ASIA',
'Haiti': 'CA',
'Poland': 'EUROPE',
'Holand-Netherlands': 'EUROPE',
'Philippines': 'ASIA',
'Vietnam': 'ASIA',
'Hungary': 'EUROPE',
'England': 'EUROPE',
'South': 'CA',
'Jamaica': 'CA',
'Honduras': 'CA',
'Portugal': 'EUROPE',
'Mexico': 'CA',
'El-Salvador': 'SA',
'India': 'ASIA',
'Puerto-Rico': 'CA',
'China': 'ASIA',
'Yugoslavia': 'EUROPE',
'United-States': 'NA',
'Trinadad&Tobago': 'CA',
'Greece': 'EUROPE',
'Japan': 'ASIA',
'Taiwan': 'ASIA',
'Nicaragua': 'CA',
'Columbia': 'SA'}

```

```

census_orig['MaritalStatus'] = census_orig['MaritalStatus'].apply(lambda
census_orig['Relationship'] = census_orig['Relationship'].apply(lambda
census_orig['WorkClass'] = census_orig['WorkClass'].apply(lambda x: wor
census_orig['Occupation'] = census_orig['Occupation'].apply(lambda x: o
census_orig['NativeCountry'] = census_orig['NativeCountry'].apply(lambda

```

```

In [27]: # Convert labels from strings to boolean
label_encoder = preprocessing.LabelEncoder()
census_labels = pd.DataFrame(label_encoder.fit_transform(census_orig.il

# Convert nominal attributes to encoded versions
attr_encoder = feature_extraction.DictVectorizer(sparse=False)
census_data = pd.DataFrame(attr_encoder.fit_transform(census_orig.iloc[
census_data.columns = attr_encoder.get_feature_names()

census_data_scaled_transformed = pd.DataFrame(preprocessing.StandardSca
#census_data_scaled.columns = attr_encoder.get_feature_names()

```

```

In [29]: # Part 3 Plotting wss
from sklearn import mixture, cluster

def wss(kmeans, cluster_size):
    wss = 0
    # Iterate over each cluster
    for cluster in range(0, cluster_size):
        cluster_instances = []
        # Iterate over each instance of the dataset
        for i in range(0, len(kmeans)):
            # If that dataset instance is in the right cluster
            if (kmeans[i] == cluster):
                cluster_instances.append(i)

        census_cluster = census_data_scaled_transformed.iloc[cluster_in
        cols = census_cluster.shape[1]

        for col in range(0, cols):
            feature_vals = census_cluster[col].values
            attr_mean = 0
            for i in range(0, len(feature_vals)):
                attr_mean += feature_vals[i]

            attr_mean = float(attr_mean / float(len(feature_vals)))

            for i in range(0, len(feature_vals)):
                point = (feature_vals[i] - attr_mean)**2
                wss+= point

    return wss

cluster_size = []
kmeans_wss_vals = []
gmm_wss_vals = []
for cluster_num in range(1, 100, 4):
    km = cluster.KMeans(n_clusters=cluster_num, init='random', random_s
    gmm = mixture.GMM(n_components=cluster_num, random_state=20160217)

    kmeans_wss = wss(km.fit_predict(census_data_scaled_transformed), cl
    gmm_wss = wss(gmm.fit_predict(census_data_scaled_transformed), clus

    cluster_size.append(cluster_num)
    kmeans_wss_vals.append(kmeans_wss)
    gmm_wss_vals.append(gmm_wss)

print "Cluster: " + str(cluster_num)
print "KM wss : " + str(kmeans_wss) + " GMM wss: " + str(gmm_wss)


```

Cluster: 1
KM wss : 723888.000004 GMM wss: 723888.000004
Cluster: 5
KM wss : 492429.705842 GMM wss: 499703.121708
Cluster: 9
KM wss : 370431.097226 GMM wss: 362805.551685
Cluster: 13
KM wss : 294261.962079 GMM wss: 259872.786937
Cluster: 17
KM wss : 220227.067138 GMM wss: 228357.23344
Cluster: 21
KM wss : 216517.091194 GMM wss: 202982.547842
Cluster: 25
KM wss : 181168.577799 GMM wss: 183007.140588
Cluster: 29
KM wss : 161262.112427 GMM wss: 175355.574745
Cluster: 33
KM wss : 152499.370529 GMM wss: 165425.2983
Cluster: 37
KM wss : 146394.944824 GMM wss: 174520.754463
Cluster: 41
KM wss : 139603.624081 GMM wss: 144054.481554
Cluster: 45
KM wss : 134380.445659 GMM wss: 141416.546179
Cluster: 49
KM wss : 125829.492452 GMM wss: 140611.220053
Cluster: 53
KM wss : 122980.486992 GMM wss: 128431.744274
Cluster: 57
KM wss : 119579.75442 GMM wss: 123476.215474
Cluster: 61
KM wss : 115743.333467 GMM wss: 124506.593976
Cluster: 65
KM wss : 111449.211269 GMM wss: 119714.445448
Cluster: 69
KM wss : 107267.157003 GMM wss: 119096.610484
Cluster: 73
KM wss : 103044.82988 GMM wss: 116943.157121
Cluster: 77
KM wss : 100190.681263 GMM wss: 110052.241624
Cluster: 81
KM wss : 98714.9060005 GMM wss: 111586.280982
Cluster: 85
KM wss : 95692.6497576 GMM wss: 114402.242063
Cluster: 89
KM wss : 93245.4641332 GMM wss: 109390.000737
Cluster: 93
KM wss : 91932.3811497 GMM wss: 105908.241896
Cluster: 97
KM wss : 90327.4640868 GMM wss: 104254.14054

```
In [30]: plt.plot(cluster_size, kmeans_wss_vals, 'b', cluster_size, gmm_wss_vals
blue_patch = mpatches.Patch(color='blue', label='kmeans wss')
green_patch = mpatches.Patch(color='green', label='gm wss')

plt.legend(handles=[blue_patch, green_patch], loc='upper left')

plt.ylabel("WSS")
plt.show()
```



```

In [31]: # Part 2 salient attribute
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

cluster_size = 40
gmm = mixture.GMM(n_components=cluster_size, random_state=20160217)

gmm_predict = gmm.fit_predict(census_data_scaled_transformed)

cluster_instances = []
# Iterate over each instance of the dataset
for i in range(0, len(gmm_predict)):
    # If that dataset instance is in the right cluster
    if (gmm_predict_pca[i] == 5):
        cluster_instances.append(i)

transformed_cluster = census_data_scaled.iloc[census_data_scaled_transformed_indices]
normal_cluster = census_data_scaled.iloc[cluster_instances]


plt.figure(figsize=(20,20))

normal_cluster.boxplot(return_type='axes')

```


Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1129634d0>






```
In [32]: plt.figure(figsize=(20,20))  
transformed_cluster.boxplot(return_type='axes')
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x109493f90>
```



```
In [ ]:
```